

Causal Aspects of Deep Reinforcement Learning

Causal Inference & Deep Learning

MIT IAP 2018

Fredrik D. Johansson




Deep reinforcement learning



Reinforcement learning in general

- Often illustrated as a loop over time $t = 0, 1, 2, \dots$




- Continuously, the agent updates its belief of the world based on the feedback from the environment
- Learning through trial and error

Maximizing reward

- ▶ The goal of most RL algorithms is to maximize the value, or expected return
- ▶ **Return:** $R = \sum_{t=1}^T r_t$ (sometimes infinite, discounted sum)
- ▶ **Value:** $V_{\pi} = \mathbb{E}[R]$ (sometimes conditioned on starting state)
- ▶ The expectation is taken with respect to scenarios acted out according to **policy** π

Great! Now let's treat patients

- ▶ Patient **state** at time S_t is like the game board
- ▶ Medical **treatments** A_t are like the actions
- ▶ **Outcomes**/progression R_t are the rewards in the game
- ▶ What could **possibly** go wrong?



1. Decision processes

2. Learning from batch (off-policy) data

3. Reinforcement learning paradigms

4. Applications

Decision processes

- ▶ The environment-agent system is called a **decision process**
- ▶ The process specifies how states S_t , actions A_t , and rewards R_t are **distributed**: $p(S_0, \dots, S_T, A_0, \dots, A_T, R_0, \dots, R_T)$
- ▶ The agent interacts with the environment according to a policy $p(A_t \mid \dots)$. (The ... depends on the type of agent)

Markov decision processes

- ▶ Markov decision processes (MDPs) are a special case

- ▶ (Unknown) Markov **transitions**:


$$p(S_t \mid S_0, \dots, S_{t-1}, A_0, \dots, A_{t-1}) = p(S_t \mid S_{t-1}, A_{t-1})$$

- ▶ (Unknown) Markov **reward** function: $p(R_t \mid S_t, A_t)$

- ▶ Markov **action** policy $p(A_t \mid S_t)$, (often denoted π or μ)


Markov assumption & MDPs

- ▶ State transitions depend only on most recent state-action
- ▶ Most common model in deep RL

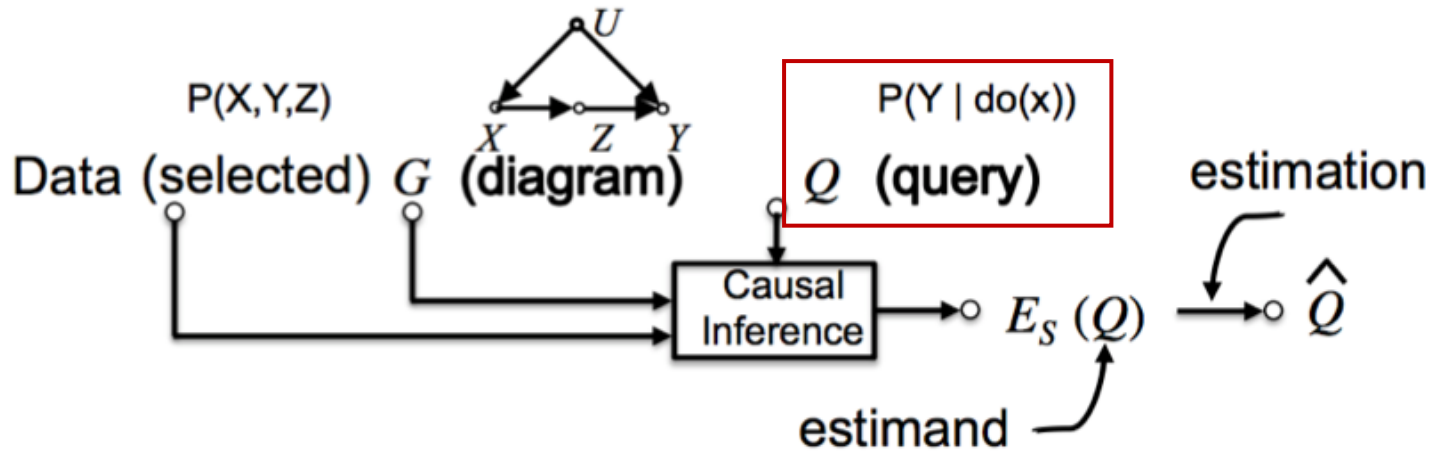


Contextual bandits

- States are independent: $p(S_t | S_{t-1}, A_{t-1}) = p(S_t)$
- Equivalent to single-step case, but focus often on exploration



What question are we asking?



What question are we asking?

- ▶ **Goals:**

- ▶ What is a policy that maximizes expected reward?
- ▶ What is the expected reward of a fixed policy π

- ▶ **Settings:**

- ▶ **On-policy:** If I can try out my new policy π in practice, how do I find the best one quickly?
- ▶ **Off-policy:** If I can't try out a policy π , how do I find a good one and evaluate it using observational (off-policy) data?

What question are we asking?

► **Goals:**

- What is a policy that maximizes expected reward?
- What is the expected reward of a fixed policy π

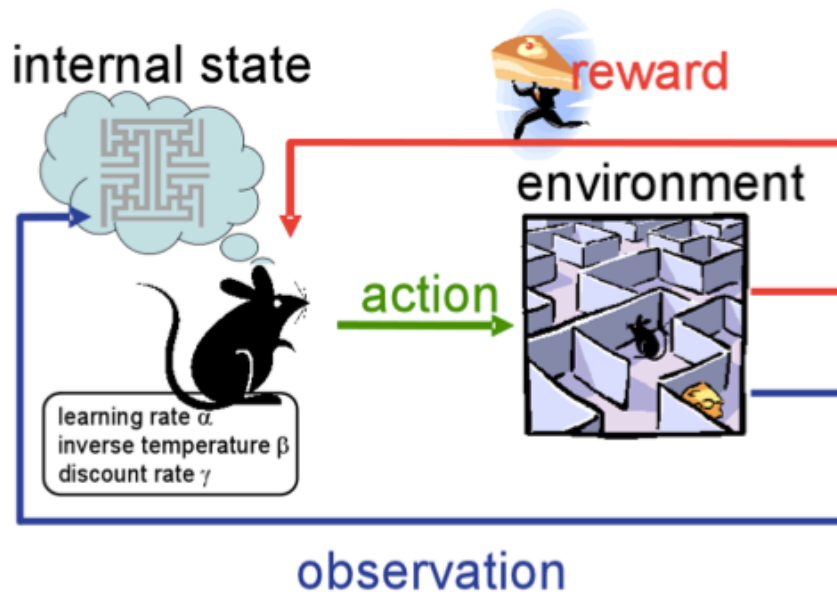
► **Settings:**

- **On-policy:** If I can try out my new policy π in practice, how do I find the best one quickly?

Focus today

- **Off-policy:** If I can't try out a policy π , how do I find a good one and evaluate it using observational (off-policy) data?

What question are we asking?



On-policy:
We are the rat.

Off-policy
We are learning from a video
of the rat in the maze.

1. Decision processes

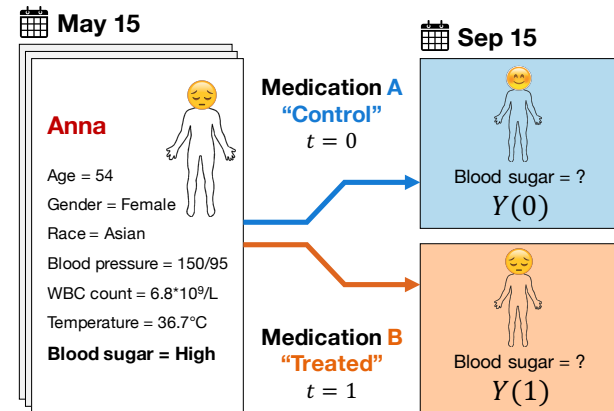
2. Learning from batch (off-policy) data

3. Reinforcement learning paradigms

4. Applications


Treating Anna over time

- ▶ Remember our diabetic patient
- ▶ We had observed hers and other patient's electronic health records over time
- ▶ Based on this information, **without experimenting** further, what would be the best treatment for Anna?



Treating Anna over time


- We assumed a simple causal graph. This let us identify the causal effect of treatment on outcome from observational data



Equivalent to a single time step MDP!

Treating Anna over time


- Let's add a time point:



Treating Anna over time

- Let's add a time point:

Anna's health status depends on how we treated her




It is likely that if Anna is diabetic, she will remain so

Treating Anna over time

- Let's add a time point:


The outcome at a later time point may depend on earlier choices



The outcome at a later time may depend on an earlier state


Treating Anna over time

- Let's add a time point:




Treating Anna over time

- Our next action and outcome may depend on the whole history



Treating Anna over time

- Not only is this a complicated causal graph, it is not a Markov decision process either!




How can we find the effect of our policy on the expected reward¹?

¹The picture is slightly misleading: which arrows we care about depend on which effect we care about

Notation

- A little necessary notation



¹The picture is slightly misleading: which arrows we care about depend on which effect we care about

Assumptions:

- Conditions for **identifiability** of potential reward:

Single-step case

Strong ignorability:

$$Y(0), Y(1) \perp\!\!\!\perp T \mid X$$

“No *hidden* confounders”

Overlap:

$$\forall x, t: p(T = t \mid X = x) > 0$$

“All actions possible”

Sequential case

Sequential randomization:

$$\mathcal{R} \perp\!\!\!\perp A_t \mid \bar{S}_t, \bar{A}_{t-1}$$

“Reward indep. of policy given history”


Positivity:

$$\forall a, t: p(A_t = a \mid \bar{S}_t, \bar{A}_{t-1}) > 0$$

“All actions possible at all times”


Summarizing history

- Conditioning on history of states and actions is algorithmically challenging: different length of history, high dimensionality etc
- Instead, we may attempt to summarize history in a variable Z



Summarizing history


- We can use sequence models such as recurrent neural networks and LSTMs to summarize state-action history
- For causal reasoning, we need assumptions to hold w.r.t. Z



?
 $\mathcal{R} \perp\!\!\!\perp A_t \mid Z_t$

Partially observable MDPs (POMDPs)

- A related concept are POMDPs, in which what we observe is a **partial/noisy version** of a latent Markov system



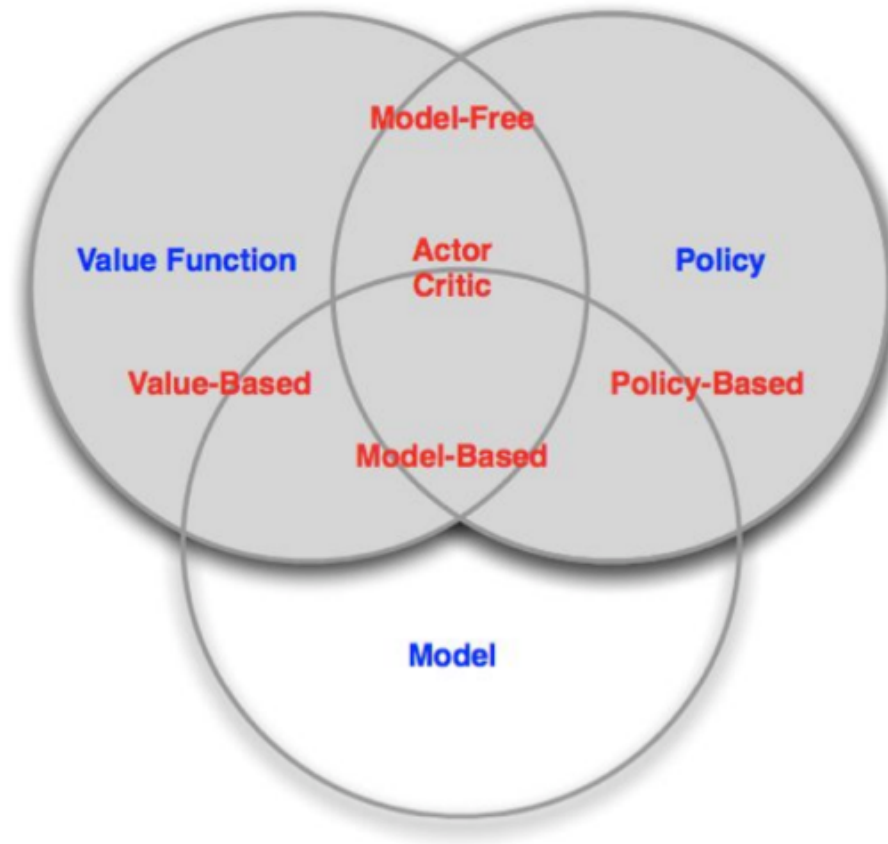
1. Decision processes

2. Learning from batch (off-policy) data

3. Reinforcement learning paradigms


4. Applications

Reinforcement learning



Model-based RL

- Explicitly model state transitions: $p(S_{t+1} \mid A_t, S_t)$
- Can be used for **planning** to discover optimal policy
- Predicts future states, and acts according to learned policy



Policy search

- ▶ Directly optimizes over the (possibly stochastic) **policy** $\pi(s_t)$,
(not using a value function)
- ▶ Can be used both to learn an observed policy (e.g. man-made)
and to search for optimal policies
- ▶ An example: Try an action out. If it had a good result, increase
the probability of that happening again.
- ▶ Typically **on-policy** (need exploration)

Model-free RL: Q-learning

- ▶ **Off-policy, value-based** reinforcement learning method
- ▶ A Q-function $Q(s, a)$ assigns a value to each state-action pair (s, a) to represent the long-term reward of that action
- ▶ The best value function equals the expected future reward of taking an action in a state

$$Q^*(s, a) = \max_{\pi} \mathbb{E}_{\pi}[R_t \mid S_t = s, A_t = a]$$

Model-free RL



Bellman equation

- ▶ The Bellman equations states that the optimal Q-function has the property (where s' is the state after taking action a in s)

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a')$$

- ▶ Q-iteration repeatedly uses this rule to update current estimate
- ▶ If the state space is finite, Q can be represented by a table, and the optimum can be found through dynamic programming

Q-learning with function approximation

- ▶ When the state space is continuous, we have to rely on function approximation of Q (as opposed to a table)
- ▶ We can still use the Bellman equation, but are no longer guaranteed to find the optimum

$$R(Q) = \mathbb{E}_{\pi} \left[\left(\underbrace{r + \gamma \max_{a'} \hat{Q}(s', a')}_{\substack{\text{Reward} \\ \text{Target,} \\ \text{typically an old estimate of } Q}} - \underbrace{Q(s, a)}_{\text{Prediction}} \right)^2 \right]$$

Q-learning with function approximation

► Typically proceeds iteratively:


1. Initialize target \hat{Q}
2. **Repeat:**
 1. Minimize Q-loss $R(Q)$ w.r.t. prediction/policy network Q
 2. After some time, update target \hat{Q} with recent Q

$$R(Q) = \mathbb{E}_{\pi} \left[\left(r + \gamma \max_{a'} \hat{Q}(s', a') - Q(s, a) \right)^2 \right]$$

Deep Q-learning

- Function approximation with deep neural networks

$$R(Q) = \mathbb{E}_{\pi} \left[\underbrace{\left(\overset{\text{Reward}}{r} + \gamma \max_{a'} \hat{Q}(s', a') \right)}_{\text{Target}} - \underbrace{Q(s, a)}_{\text{Prediction}} \right]^2$$



Same architecture, different weights


Q-learning with function approximation

- ▶ Optimization dynamics:
 - ▶ The goal post \hat{Q} (target) keeps changing as we update Q
 - ▶ No guarantee to converge to optimal Q in general case
- ▶ To be causally sound in non-Markov case, we should predict Q from whole history. This is not typically done!
- ▶ Distributional shift (like in single-step case):
 - ▶ We don't have samples from the policy we are evaluating!
 - ▶ Our loss function is an expectation under a distribution from which we have no samples!

Counterfactuals are very different!


► **Single-step case:**

Counterfactuals are changes to a single action
(control->treated and vice versa)



► **Sequential case:**

Counterfactuals are changes to whole sequences of actions!



Coping with distributional shift Pt. II

- Last time we discussed importance sampling estimators

$$\mathbb{E}_{\mathbf{q}(\mathbf{x})}[f(\mathbf{x})] = \mathbb{E}_{\mathbf{p}(\mathbf{x})}\left[f(\mathbf{x}) \frac{q(\mathbf{x})}{p(\mathbf{x})}\right] \approx \frac{1}{n} \sum_{i=1}^n \frac{q(x_i)}{p(x_i)} f(x_i)$$

- Where $x_1, \dots, x_n \sim p(x)$
- What happens when x is a sequence?

Importance sampling for RL

- ▶ First of all, we don't observe i.i.d. samples of state-action transitions, but whole sequences
- ▶ We let the loss be the expected total loss over a sequence

$$R(Q) = \mathbb{E}_{\pi} \left[\sum_{t=1}^T \underbrace{\left(r_t + \gamma \max_{a'} \hat{Q}(s_{t+1}, a') - Q(s_t, a_t) \right)^2}_{L_t = L(s_t, a_t, s_{t+1})} \right]$$

Importance sampling for RL

- Product of importance weights over time

$$\mathbb{E}_{\mathbf{q}} \left[\sum_{t=1}^T L_t \right] = \mathbb{E}_{\mathbf{p}} \left[\sum_{t=1}^T L_t \prod_{t'=0}^t \frac{\mathbf{q}(a_{t'} | s_{t'})}{\mathbf{p}(a_{t'} | s_{t'})} \right]$$

- Could have extremely high variance if **product** of large factors (if **p** and **q** are very different)
- Effective sample size $ESS = \frac{(\sum w_i)^2}{\sum w_i^2}$, where w_i weight of sequence

Trading off bias for variance

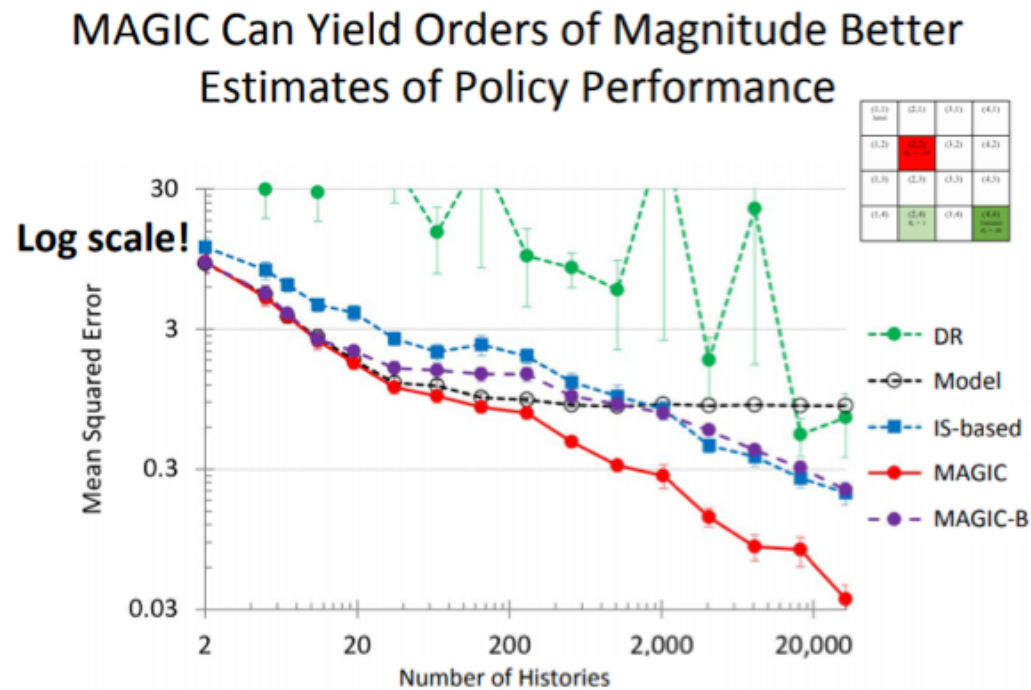
- ▶ Even if importance sampling yields an unbiased estimator of the policy value, a biased version with smaller variance might be preferred
- ▶ Think of bias-variance decomposition:

$$MSE = Bias + Variance + Noise$$

- ▶ Can get smaller prediction error by trading off bias and variance

MAGIC¹

- Combining model-based and model-free RL with IS



¹Thomas and Brunskill, *ICML*, 2016

1. Decision processes
2. Learning from batch (off-policy) data
3. Reinforcement learning paradigms
- 4. Applications**

What made success possible/easier?

- ▶ **Full observability**

Everything important to optimal action is observed

- ▶ **Markov** dynamics

History is unimportant given recent state(s)



- ▶ Limitless **exploration** & self-play through simulation

We can test “any” policy and observe the outcome

- ▶ **Noise-less** state/outcome (for games, specifically)

Many concerns

- ▶ **Can we summarize history well?**

If we measured everything we are theoretically OK. But did we keep everything we need in our summary?

- ▶ **Is overlap/positivity enough?**

Even if it gives us unbiased estimators, what is the sample complexity? Is it reasonable?

- ▶ **Do we have the right reward function?**

Many concerns

- ▶ **Hidden confounding**

Did we measure all the necessary variables?

- ▶ **Expectation vs risk**


What can we tolerate in terms of outlier behavior?

- ▶ **Evaluation**

Can we trust our models estimates?

Sepsis treatment¹

- ▶ Work on using Q-learning to discover the right treatment for sepsis
- ▶ Dosage of a) fluids, and b) vasopressors
- ▶ Compare found policy to physician's in terms of mortality



¹Raghu et al. *MLHC*, 2017

Conclusions

- ▶ Off-policy reinforcement learning is **strictly** harder than counterfactual estimation
- ▶ The causal problems with standard regression are even greater
- ▶ Both conceptual/theoretical and practical challenges remain
- ▶ We need to take care that we are trying to solve a problem that is actually interesting